K-means clustering: Takeaways ⤴

Syntax

- Computing the Euclidean distance in Python:

```python
def calculate_distance(vec1, vec2):
    root_distance = 0
    for x in range(0, len(vec1)):
        difference = centroid[x] - player_values[x]
        squared_difference = difference**2
        root_distance += squared_difference
    euclid_distance = math.sqrt(root_distance)
    return euclid_distance
```

- Assigning observations to clusters:

```python
def assign_to_cluster(row):
    lowest_distance = -1
    closest_cluster = -1
    for cluster_id, centroid in centroids_dict.items():
        df_row = [row['ppg'], row['atr']]
        euclidean_distance = calculate_distance(centroid, df_row)
        if lowest_distance == -1:
            lowest_distance = euclidean_distance
            closest_cluster = cluster_id
        elif euclidean_distance < lowest_distance:
            lowest_distance = euclidean_distance
            closest_cluster = cluster_id
    return closest_cluster
```

- Initializing the KMeans class from scikit-learn:

```python
from sklearn.cluster import KMeans
kmeans_model = KMeans(n_clusters=2, random_state=1)
```

Concepts

- Centroid-based clustering works well when the clusters resemble circles with centers.
- K-Means clustering is a popular centroid-based clustering algorithm. The K refers to the number of clusters we want to segment our data into. K-Means clustering is an iterative algorithm that switches between recalculating the centroid of each cluster and the items that belong to each cluster.
- Euclidean distance is the most common technique used in data science for measuring distance between vectors. The formula for distance in two dimensions is:

$$\sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \ldots + (q_n - p_n)^2}$$

  where q and p are the two vectors we are comparing.
  - Example: If `q` is (5,2) and `p` is (3,1), the distance comes out to:
  $$\sqrt{(5 - 3)^2 + (2 - 1)^2} = \sqrt{5} \approx 2.23607$$

- If clusters look like they don't move a lot after every iteration, this means two things:
  - K-Means clustering doesn't cause massive changes in the makeup of clusters between iterations, meaning that it will always converge and become stable.
  - Where we pick the initial centroids and how we assign elements to clusters initially matters a lot because K-Means clustering is conservative between iterations.

- To counteract the problems listed above, the `sklearn` implementation of K-Means clustering does some intelligent things like re-running the entire clustering process lots of times with random initial centroids so the final results are a little less biased.

Resources

- Sklearn implementation of K-Means clustering
- Implementing K-Means clustering from scratch